

Efficient Neural Architecture Search via Parameter Sharing



Pattern Recognition & Machine Learning Laboratory

Sangho Kim

Aug 12, 2021



Introduction

Automated Machine Learning (AutoML)

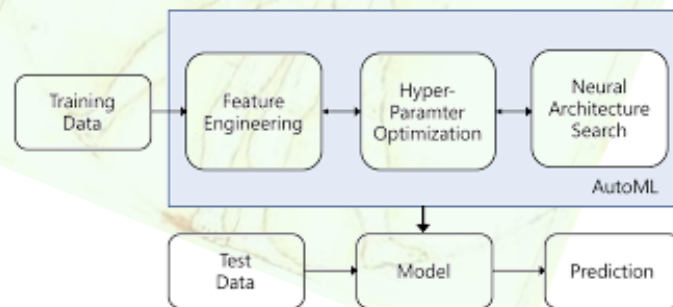
➤ The process of automating the process of applying machine learning

- Feature engineering
- Hyper-parameter optimization
- Neural Architecture Search (NAS)


Neural Architecture Search

➤ Neural Architecture Search with Reinforcement Learning [Zoph et al., 2017]

- The controller of the NAS uses the RNN to make sample architecture
- Learning the child network to extract accuracy for the validation set
- Using validation accuracy like a reward for reinforcement learning
- Learning the controller to maximize accuracy



Representative fields of AutoML



Simple NAS structure



Efficient Neural Architecture Search via Parameter Sharing [H. Pham, et al., 2018]

■ Goal


- Improve the computational complexity of NAS
- Improve the processing time of NAS

■ Motivation

- NAS uses 450 GPUs for 4 days
- Computational bottleneck of NAS is the training of each child model
- Throwing away all the trained weights
- NAS's search space can be represented in directed acyclic graph (DAG)
 - Nodes are local computation
 - Edges are information flow
 - The whole DAG means entire search space
 - The red arrows are sub-graph

■ Contributions

- Efficient Neural Architecture Search (ENAS) forces all child models to share weights
- Using single GPU, the search for architectures takes less than 16 hours
- Compared to NAS, time reduction is more than 1000 times



Directed acyclic graph



Methods (1/4)

■ Designing recurrent cells

➤ The controller RNN samples two decisions

- Which edges are activated
- Which activation functions are performed at each node

➤ Mechanism via 4 nodes

- Node 1 : $k_1 = \tanh(x_t \cdot W^{(x)} + h_{t-1} \cdot W_1^{(h)})$
- Node 2 : $k_2 = \text{ReLU}(k_1 \cdot W_{2,1}^{(h)})$
- Node 3 : $k_3 = \text{ReLU}(k_2 \cdot W_{3,2}^{(h)})$
- Node 4 : $k_4 = \tanh(k_1 \cdot W_{4,1}^{(h)})$
- Output : $h_t = (k_3 + k_4) / 2$

➤ All recurrent cells share the same set of parameters

➤ 4 activation functions are allowed


- tanh, ReLU, identity, sigmoid

➤ Search space


- Recurrent cells have N nodes
 - $4^N \times (N - 1)!$
 - If $N = 12$, there are approximately 10^{14} models in the search space

- k_ℓ : node ℓ of the cell computation
- x_t : input signal
- h_{t-1} : previous step output
- $W_{\ell,j}^{(h)}$: parameter matrix

Variable explanation



The DAG of recurrent cells





The recurrent cells



Methods (2/4)

■ Designing convolutional networks

- The controller RNN samples two decisions
 - Which edges are activated
 - Which computation operations are performed at each node
- Previous connected nodes are allowed to form skip connections
- Mechanism via 4 layers
 - Layer 2 concatenates the outputs of layer 1
 - Layer 3 concatenates the outputs of layer 1 and 2
 - Layer 4 concatenates the outputs of layer 1 and 3
- 6 operations are allowed
 - Convolutions with filter sizes 3×3 and 5×5
 - Depthwise-separable convolutions with filter sizes 3×3 and 5×5
 - Max pooling and average pooling of kernel size 3×3
- Search space
 - Convolutional networks have L layers
 - $6^L \times 2^{L(L-1)/2}$
 - If $L = 12$, there are approximately 1.6×10^{29} possible networks





Methods (3/4)

■ Designing convolutional cells

➤ Design smaller modules and connect them together to form a network


- Connect N convolution cells and 1 reduction cell

➤ The controller RNN samples two decisions


- Which edges are activated
- Which computation operations are performed at each node

➤ Mechanism via 4 nodes


- Nodes 1, 2
 - Input nodes
- Node 3
 - Sample node 2, node 2
 - Sample separable_conv 5×5 , identity
- Node 4
 - Sample node 3, node 1
 - Sample avg_pool 3×3 , separable_conv 3×3
- Output



Final network of convolutional cells



The DAG of convolutional cells



The convolutional cell



Methods (4/4)

■ Designing convolutional cells

➤ Design reduction cells

- Sample a computational graph from the search space
- Apply all operations with a stride of 2
- Reduce the spatial dimensions by a factor of 2

➤ 5 operations are allowed

- Identity
- Separable convolutions with kernel sizes 3×3 and 5×5
- Max pooling and average pooling of kernel size 3×3

➤ Search space

- Convolutional cells have B nodes
 - $(5 \times (B - 2)!)^4$
 - If $B = 7$, there are approximately 1.3×10^{11} possible networks




Experiments (1/3)

Recurrent cells on the Penn Treebank

- Running on a single Nvidia GTX 1080Ti GPU
- ENAS finds a recurrent cell in about 10 hours
- Achieves a test perplexity of 56.3, which is on par with SOTA of 56.0
- ENAS outperforms NAS by more than 6 perplexity points
- ENAS is more than 1000 times faster than NAS in terms of GPU time

| Architecture | Additional Techniques | Params (million) | Test PPL |
|-------------------------------|-----------------------------|------------------|-------------|
| LSTM (Zaremba et al., 2014) | Vanilla Dropout | 66 | 78.4 |
| LSTM (Gal & Ghahramani, 2016) | VD | 66 | 75.2 |
| LSTM (Inan et al., 2017) | VD, WT | 51 | 68.5 |
| RHN (Zilly et al., 2017) | VD, WT | 24 | 66.0 |
| LSTM (Melis et al., 2017) | Hyper-parameters Search | 24 | 59.5 |
| LSTM (Yang et al., 2018) | VD, WT, ℓ_2 , AWD, MoC | 22 | 57.6 |
| LSTM (Merity et al., 2017) | VD, WT, ℓ_2 , AWD | 24 | 57.3 |
| LSTM (Yang et al., 2018) | VD, WT, ℓ_2 , AWD, MoS | 22 | 56.0 |
| NAS (Zoph & Le, 2017) | VD, WT | 54 | 62.4 |
| ENAS | VD, WT, ℓ_2 | 24 | 56.3 |

The perplexity on Penn Treebank of ENAS



The RNN cell ENAS discovered for Penn Treebank



Experiments (2/3)

Convolutional architecture on the CIFAR-10

Entire convolutional networks

- Increasing the number of filters in ENAS shows an error rate of 3.87%, not far from the error rate of the NAS's highest model, 3.65%
- ENAS reduce the number of GPU-hours by more than 50,000 times compared to NAS.

Convolutional cells

- ENAS with cutout achieves to 2.89% test error, on par with the 2.65% by NASNet-A with cutout

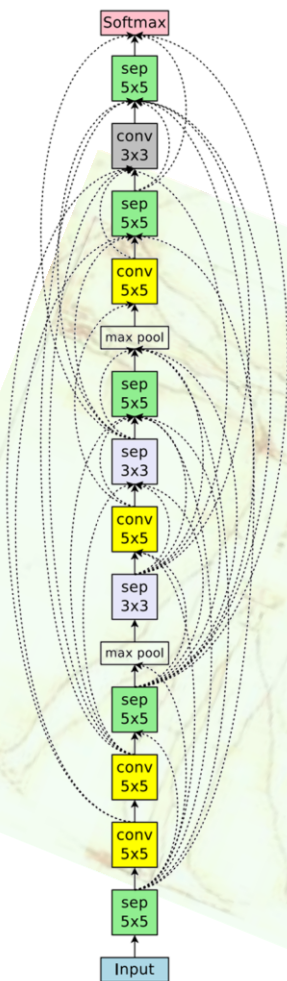
| Method | GPUs | Times (days) | Params (million) | Error (%) |
|--|------|--------------|------------------|-------------|
| DenseNet-BC (Huang et al., 2016) | — | — | 25.6 | 3.46 |
| DenseNet + Shake-Shake (Gastaldi, 2016) | — | — | 26.2 | 2.86 |
| DenseNet + CutOut (DeVries & Taylor, 2017) | — | — | 26.2 | 2.56 |
| Budgeted Super Nets (Veniat & Denoyer, 2017) | — | — | — | 9.21 |
| ConvFabrics (Saxena & Verbeek, 2016) | — | — | 21.2 | 7.43 |
| Macro NAS + Q-Learning (Baker et al., 2017a) | 10 | 8-10 | 11.2 | 6.92 |
| Net Transformation (Cai et al., 2018) | 5 | 2 | 19.7 | 5.70 |
| FractalNet (Larsson et al., 2017) | — | — | 38.6 | 4.60 |
| SMASH (Brock et al., 2018) | 1 | 1.5 | 16.0 | 4.03 |
| NAS (Zoph & Le, 2017) | 800 | 21-28 | 7.1 | 4.47 |
| NAS + more filters (Zoph & Le, 2017) | 800 | 21-28 | 37.4 | 3.65 |
| ENAS + macro search space | 1 | 0.32 | 21.3 | 4.23 |
| ENAS + macro search space + more channels | 1 | 0.32 | 38.0 | 3.87 |
| Hierarchical NAS (Liu et al., 2018) | 200 | 1.5 | 61.3 | 3.63 |
| Micro NAS + Q-Learning (Zhong et al., 2018) | 32 | 3 | — | 3.60 |
| Progressive NAS (Liu et al., 2017) | 100 | 1.5 | 3.2 | 3.63 |
| NASNet-A (Zoph et al., 2018) | 450 | 3-4 | 3.3 | 3.41 |
| NASNet-A + CutOut (Zoph et al., 2018) | 450 | 3-4 | 3.3 | 2.65 |
| ENAS + micro search space | 1 | 0.45 | 4.6 | 3.54 |
| ENAS + micro search space + CutOut | 1 | 0.45 | 4.6 | 2.89 |

Classification errors of ENAS and baselines on CIFAR-10





Experiments (3/3)

Convolutional architecture on the CIFAR-10



The convolutional network ENAS discovered for CIFAR-10



The convolutional cells ENAS discovered for CIFAR-10